PGConf.BE 2026

contentgrid BY AMEXIO

Scaling Semantic Models

Bespoke PostgreSQL Schemas at SaaS Scale

Thijs Lemmens · May 2026

About Xenit & Amexio Group

Alfresco ECM Service Provider

Long-time Alfresco implementation partner — now building and offering ContentGrid as our own cloud-native ECM product.

European Digital Solutions Group

Pan-European group specialising in ECM, digital transformation, and intelligent information management.

Xenit was acquired by Amexio Group in April 2025 — bringing ContentGrid into a broader portfolio of ECM expertise across Europe.

Agenda

1. The problem with EAV / generic data models
2. Intro to ContentGrid
3. Attribute based access controll
4. Roadmap
5. Conclusions

contentgrid
BY AMEXIO

The Problem

Why generic EAV tables fall short

contentgrid
BY AMEXIO

The EAV Trap

How most CMS platforms store content

Generic EAV schema

```
CREATE TABLE alf_node (
  id      bigint PRIMARY KEY,
  type_qname_id bigint REFERENCES alf_qname
);
CREATE TABLE alf_qname (
  id        bigint PRIMARY KEY,
  local_name varchar(200)
);
CREATE TABLE alf_node_properties (
  node_id       bigint REFERENCES alf_node,
  qname_id      bigint REFERENCES alf_qname,
  string_value  varchar(1024),
  boolean_value bit,
  int_value     int8,
  double_value  float8,
  ...
);
```

What you actually want

```
CREATE TABLE articles (
  id           uuid PRIMARY KEY,
  title        text        NOT NULL,
  published_at timestamptz,
  word_count   integer,
  author_id    uuid REFERENCES authors
);
```

contentgrid BY AMEXIO

One Document in EAV

Storing a single article takes rows across 3 tables

alf_node – 1 row

ID	TYPE_QNAME_ID
42	7

alf_qname – excerpt

ID	LOCAL_NAME
1	title
2	published_at
3	word_count
4	author_id
7	article

alf_node_properties – 4 rows (the actual values)

NODE_ID	QNAME_ID	STRING_VALUE	BOOLEAN_VALUE	INT_VALUE	DOUBLE_VALUE
42	1	**Getting Started with PostgreSQL**	*NULL*	*NULL*	*NULL*
42	2	**2024-03-15T09:00:00Z**	*NULL*	*NULL*	*NULL*
42	3	*NULL*	*NULL*	**1842**	*NULL*
42	4	**usr_abc123**	*NULL*	*NULL*	*NULL*

Same Document, Bespoke Table

One row. Every column typed and named.

articles — 1 row

id uuid	title text	published_at timestamptz	word_count integer	author_id uuid
42	Getting Started with PostgreSQL	2024-03-15 09:00:00+00	1842	usr_abc123

1 row · 0 NULLs · every value in the right type

contentgrid BY AMEXIO

EAV: The Hidden Costs (1/2)

Data integrity

Concern	EAV	Native schema
Type safety	One column per data type	Enforced by the column type
Constraints	`NOT NULL`, `UNIQUE` cannot be expressed	Declarative, DB-enforced
Referential integrity	Application-enforced	Foreign keys

EAV: The Hidden Costs (2/2)

Performance & tooling

Concern	EAV	Native schema
Query complexity	Reconstruct one node via dozens of `JOIN`s	Plain `SELECT`
Query planning	Value-column statistics mix all property types — per-attribute selectivity is invisible to the planner	Accurate per-column statistics
Index efficiency	Index on `(qname_id, string_value)` covers all attributes of type string	Targeted per-column indexes
Tooling compatibility	Breaks ORMs, analytics tools	Works out of the box
Search performance	Sync to Solr/Elasticsearch — operational heavy, denormalizes relations, eventual consistency	Native full-text search, joins, transactional consistency

> *EAV trades correctness and performance for schema flexibility — but you can have both.*

How Bad Is It? A Real Migration

The scenario

- **250 million** documents
- **60 attributes** each
- EAV model on Oracle

How large is the Oracle database?

A	< 500 GB
B	500 GB – 2 TB
C	2 TB – 5 TB
D	> 5 TB

contentgrid BY AMEXIO

The EAV Oracle database: 6 TB

- 15 billion property rows, each with 6+ typed value columns
- Indexes on `(qname_id, *_value)` for every attribute type

> *Now — same data, native PostgreSQL schema. Your guess?*

contentgrid
BY AMEXIO

Same Data, Native PostgreSQL Schema

What changes

- One row per document
- Typed columns — no attribute explosion

How large is the PostgreSQL database?

A	Still ~6 TB
B	1 TB – 3 TB
C	250 GB – 1 TB
D	< 250 GB

contentgrid
BY AMEXIO

The ContentGrid database: ~300 GB

- 250M rows of typed data
- Targeted indexes
- No attribute row explosion

20× smaller — by fixing the data model

contentgrid
BY AMEXIO

Storage at a Glance

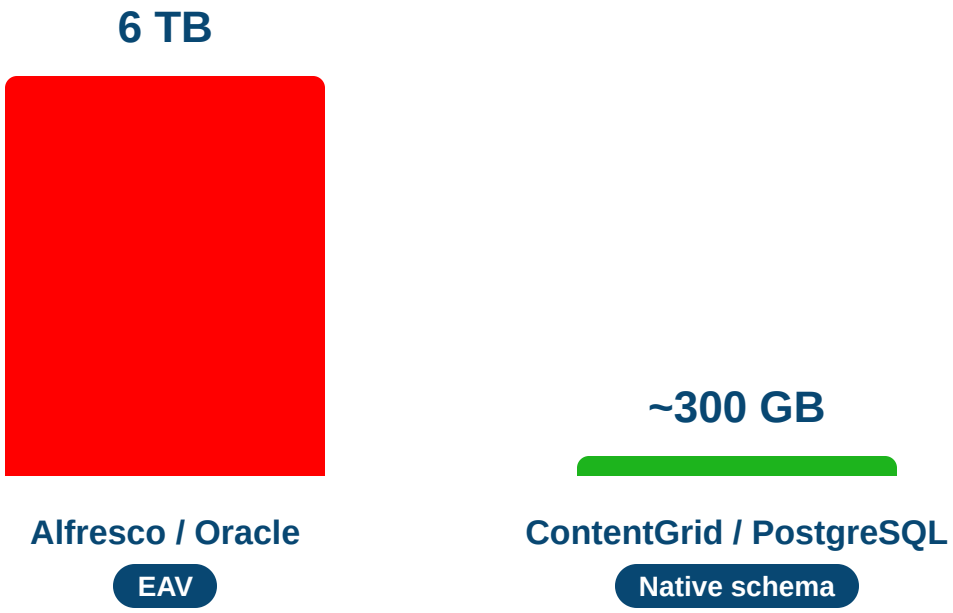

Our Approach

Generating native schemas from semantic models

contentgrid
BY AMEXIO

Two Platforms

One to define the model. One to run it.

MANAGEMENT PLATFORM

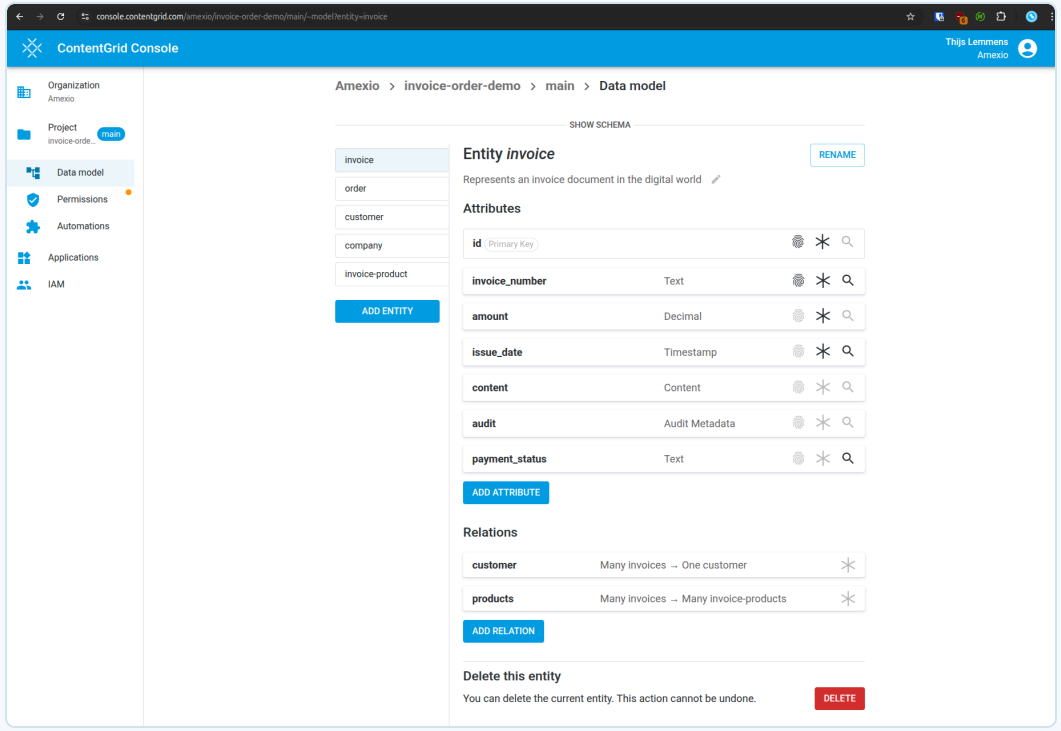

ContentGrid Console — define entities, attributes, permissions

RUNTIME PLATFORM

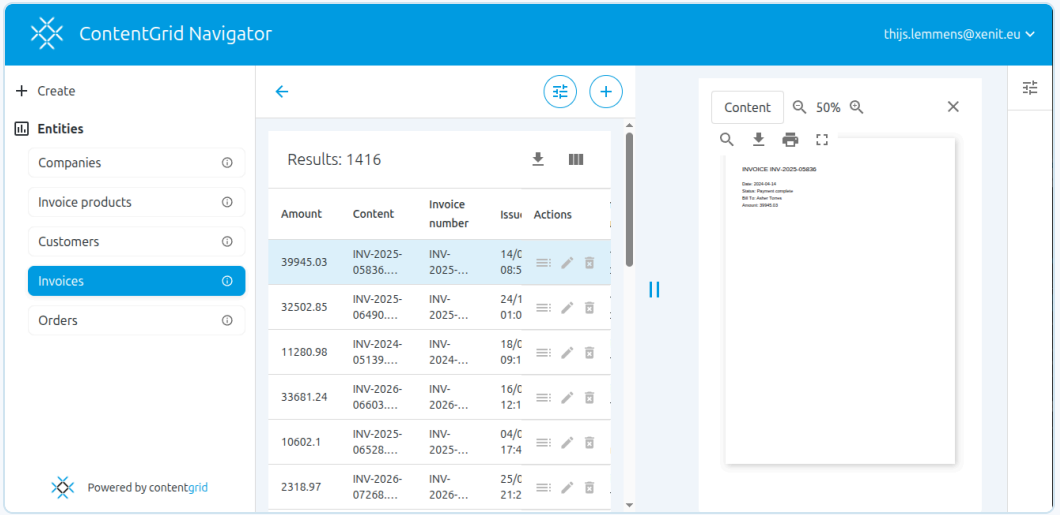

ContentGrid Navigator — end-user app, driven by the model

contentgrid BY AMEXIO

The Management Platform

Architect

Source of truth for the domain model — entities, attributes, relations

→

Scribe

Compiles model changes into:

- JSON model
- Flyway SQL migrations
- OPA policies

→

Captain

- Provisions database
- manages credentials
- deployment

Every model change → versioned, reviewed SQL migration · Schema always reflects the model — no manual DDL

contentgrid
BY AMEXIO

Model-Driven Schema Generation

User-defined semantic model

```
{
  "name": "article",
  "table": "article",
  "attributes": [
    {
      "type": "simple",
      "name": "title",
      "dataType": "text",
      "columnName": "title",
      "constraints": [{ "type": "required" }]
    },
    {
      "type": "simple",
      "name": "publishedAt",
      "dataType": "datetime",
      "columnName": "published_at"
    },
    {
      "type": "simple",
      "name": "wordCount",
      "dataType": "long",
      "columnName": "word_count"
    }
  ],
  ...
}
```

Generated PostgreSQL DDL

```
CREATE TABLE article (
  id           uuid PRIMARY KEY
                 DEFAULT gen_random_uuid(),
  title        text NOT NULL,
  published_at timestamptz,
  word_count   bigint
);
```

contentgrid BY AMEXIO

The Runtime Platform

How a request flows from client to database

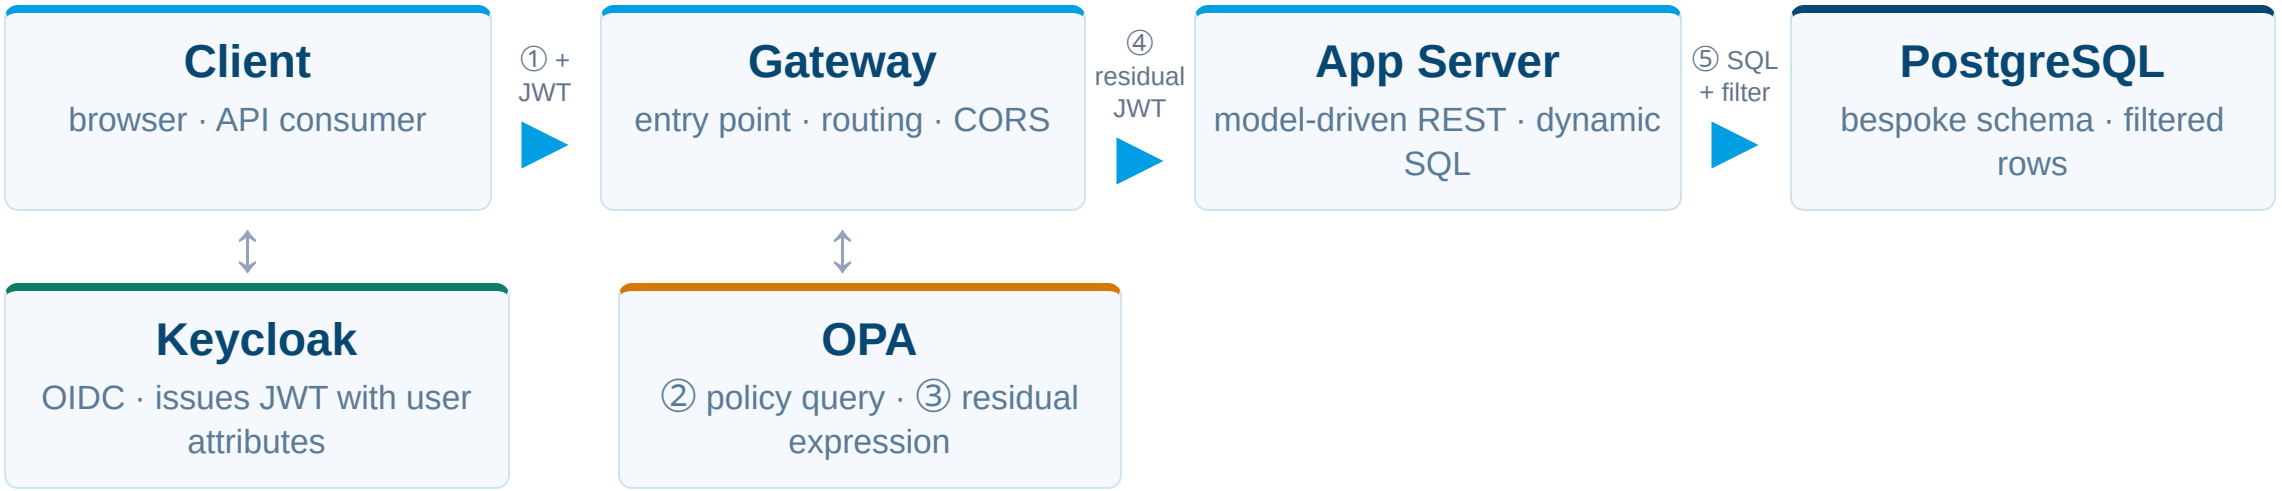

Gateway encodes OPA's residual as a JWT claim — App Server applies it as a SQL `WHERE` filter

contentgrid BY AMEXIO

Partial Evaluation: The Key Insight

Naïve approach

```
for each row in invoices:        ← full table scan
  if OPA.check(row, user):       ← N network calls
    return row
```

N rows fetched, N OPA calls, unauthorized data in memory

Partial evaluation

```
OPA.partial_eval(policy, user)  ← 1 call, no entity data
    → residual: dept='sales'
      OR status='published'

SELECT * FROM invoices
WHERE dept='sales'
  OR status='published'          ← filtered at the DB
```

1 OPA call, unauthorized rows never leave PostgreSQL

The security boundary is enforced inside the database — not in application memory

contentgrid BY AMEXIO

From Policy Rule to SQL Filter

① Policy condition

```
entity.department == user.department
OR entity.status == "published"
```

② OPA partial eval

```
OPA.partial_eval(policy, user)
  → deptartment == "sales"  OR  status == "published"
```

③ SQL WHERE (JOOQ)

```
SELECT * FROM invoices
WHERE department = 'sales' OR status = 'published'
```

Defined in Architect UI → compiled to Rego by Scribe → evaluated by OPA (1 call, result in Gateway JWT) → App Server builds JOOQ predicate

Unauthorized rows never leave the database — the filter runs inside PostgreSQL

contentgrid BY AMEXIO

Roadmap

- Full text search
- Facetted Search
- Searching over multiple entities (tables)
- Versioning (QDA)

contentgrid
BY AMEXIO

Conclusions

- EAV trades correctness & performance for flexibility — you can have both
- Native schemas: 20× smaller, full SQL tooling, real constraints
- ContentGrid generates & migrates schemas from a semantic model
- Security boundary enforced inside PostgreSQL via partial evaluation

contentgrid
BY AMEXIO

Questions?

contentgrid.com

contentgrid
BY AMEXIO

Thank You

Thijs Lemmens

thijs.lemmens@xenit.eu